

Understanding the Poisson solution

CG, preconditioning, and performance

Ed Bueler

Dept. of Mathematics and Statistics, UAF

23 February 2016

outline for today

finish Chapter 3 of book:

- ▶ recall Poisson code
- ▶ convergence
- ▶ CG = conjugate gradients
- ▶ preconditioning
- ▶ scaling flaw in preconditioned CG

next week:

- ▶ start Chapter 4 on Newton's method

recall c/ch3/poisson.c

- ▶ `poisson.c` discretizes on a structured grid:

$$-\nabla^2 u = f \quad \rightarrow \quad \mathbf{A} \mathbf{u} = \mathbf{b}$$

- ▶ it then hands the linear system $\mathbf{A} \mathbf{u} = \mathbf{b}$ to PETSc:

```
formRHS(da,b);  
formMatrix(da,A);  
KSPCreate(PETSC_COMM_WORLD,&ksp);  
KSPSetOperators(ksp,A,A);  
KSPSetFromOptions(ksp);  
KSPSolve(ksp,b,u);
```


goals once you have an initial implementation

- ▶ **convergence**: is numerical method implemented correctly?
- ▶ **exposure**: what is occurring inside PETSc when we solve?
- ▶ **efficiency**: how to get high performance?

evidence for convergence

```
$ for K in 0 1 2 3 4 5 6; do ./poisson -da_refine $K; done
on 9 x 9 grid: error |u-uexact|_inf = 0.000763959
on 17 x 17 grid: error |u-uexact|_inf = 0.000196764
on 33 x 33 grid: error |u-uexact|_inf = 4.91557e-05
on 65 x 65 grid: error |u-uexact|_inf = 1.29719e-05
on 129 x 129 grid: error |u-uexact|_inf = 3.76924e-06
on 257 x 257 grid: error |u-uexact|_inf = 1.73086e-06
on 513 x 513 grid: error |u-uexact|_inf = 1.23567e-06
```

- ▶ results above are ★ on graph
- ▶ add `-ksp_rtol 1.0e-12` to get ● results
- ▶ or `-ksp_type preonly`
`-pc_type lu`
- ▶ FD method can give $O(h^2)$ at best, so `poisson.c` is *correct*



ask PETSc to show parallel structured grid

- ▶ use option `-dm_view`:


```
$ ./poisson -dm_view
```

```
$ mpiexec -n 2 ./poisson -dm_view
```

- ▶ with `draw`:

```
$ mpiexec -n 2 ./poisson -dm_view draw -draw_pause 2
```

```
$ mpiexec -n 4 ./poisson -da_grid_x 5 -da_grid_y 7 \  
-dm_view draw -draw_pause 2
```



matrix structure in poisson.c

- ▶ **note line** `MatSetOptionsPrefix(A, "a_")` in `poisson.c`

- ▶ **use** `-a_mat_view`:

```
$ ./poisson -da_grid_x 4 -da_grid_y 4 -a_mat_view  
$ ./poisson -da_grid_x 4 -da_grid_y 4 \  
    -a_mat_view ::ascii_dense  
$ ./poisson -a_mat_view draw -draw_pause 5
```

- ▶ **use it to read matrix into MATLAB/OCTAVE:**

```
$ ./poisson -da_grid_x 5 -da_grid_y 7 \  
    -a_mat_view ascii:foo.m:ascii_matlab  
$ octave  
  >> foo  
  >> whos  
  >> spy(...)
```

try various KSP and PC

- ▶ defaults include `gmres`, `ilu`
- ▶ matrix is symmetric so `minres` also available
- ▶ matrix is positive-definite so `cg`, `icc`, `cholesky` also
- ▶ table below from

```
$ timer ./poisson -da_refine 5 \  
    -ksp_converged_reason \  
    -ksp_type KSP -pc_type PC
```

<u>KSP</u>	<u>PC</u>	<u>time (s)</u>	<u>iterations</u>
gmres	none	8.44	4705
	ilu	1.35	506
	ilu + restart=200	1.46	174
cg	none	0.52	606
	jacobi	0.57	606
	icc	0.33	177
	icc + rtol= 10^{-14}	0.38	261
preonly	cholesky	8.66	1
minres	none	0.76	579

Table 3.1: Times and number of KSP iterations for serial runs of `poisson.c` on 257×257 grids. The assembled matrix is *symmetric*, *diagonally-dominant*, and *positive definite*. All runs were on WORKSTATION (see page 43).

CG = conjugate gradients

- ▶ `-ksp_type cg`
- ▶ well-known Krylov iteration for solving $\mathbf{A}\mathbf{u} = \mathbf{b} \rightarrow$
- ▶ requires \mathbf{A} to be symmetric and positive-definite
 - sometimes converges even if not
 - reliable compared to Richardson!
- ▶ less expensive than GMRES
 - 4 vectors in memory vs many (~ 30)
 - note: both iterations do one product $\mathbf{A}\mathbf{v}$ per iteration

$$\mathbf{r}_0 := \mathbf{b} - \mathbf{A}\mathbf{x}_0$$

$$\mathbf{p}_0 := \mathbf{r}_0$$

$$k := 0$$

repeat

$$\alpha_k := \frac{\mathbf{r}_k^\top \mathbf{r}_k}{\mathbf{p}_k^\top \mathbf{A} \mathbf{p}_k}$$

$$\mathbf{x}_{k+1} := \mathbf{x}_k + \alpha_k \mathbf{p}_k$$

$$\mathbf{r}_{k+1} := \mathbf{r}_k - \alpha_k \mathbf{A} \mathbf{p}_k$$

if \mathbf{r}_{k+1} is sufficiently small then exit

$$\beta_k := \frac{\mathbf{r}_{k+1}^\top \mathbf{r}_{k+1}}{\mathbf{r}_k^\top \mathbf{r}_k}$$

$$\mathbf{p}_{k+1} := \mathbf{r}_{k+1} + \beta_k \mathbf{p}_k$$

$$k := k + 1$$

end repeat

The result is \mathbf{x}_{k+1}

what CG accomplishes

- ▶ let \mathbf{u}_j be iterates from (un-preconditioned) CG
- ▶ errors $\mathbf{e}_j = \mathbf{u}_j - \mathbf{u}$ satisfy

$$\|\mathbf{e}_j\|_A = \min_{p \in \mathcal{P}_j^1} \|p(A)\mathbf{e}_0\|_A$$

- where $\mathcal{P}_j^1 = \{\text{degree } j \text{ polynomials } p(x) \text{ with } p(0) = 1\}$
 - $\|\mathbf{v}\|_A = \langle \mathbf{v}, A\mathbf{v} \rangle^{1/2}$... not the L^2 norm
 - thus $\|\mathbf{e}_j\|_A$ decrease monotonically
- ▶ the A -norm of the error is controlled by $\kappa = \kappa_2(A)$:

$$\frac{\|\mathbf{e}_j\|_A}{\|\mathbf{e}_0\|_A} \leq 2 \left(\frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1} \right)^j$$

preconditioning CG

- ▶ to improve CG performance, preconditioning the system

$$A\mathbf{u} = \mathbf{b} \quad \rightarrow \quad M^{-1}A\mathbf{u} = M^{-1}\mathbf{b}$$

- ▶ last slide says: preconditioning needs to reduce the condition number
- ▶ we need M symmetric and positive-definite, and factored:

$$M = EE^T$$

- ▶ goal:

$$\kappa_2(A) \gg \kappa_2(E^{-1}AE^{-T})$$

- ▶ “incomplete cholesky” (`-pc_type icc`) computes E which is a sparse approximate factor of A :

$$EE^T \approx A$$

- ▶ `-pc_type cholesky` does it perfectly: $EE^T = A$

measuring performance of `poisson.c`

- ▶ timing should use `-with-debugging=0` PETSc configuration
- ▶ and we should use a fine grid so that timing is less dominated by other system activities

```
$ export PETSC_ARCH=linux-c-opt
$ make poisson
...
$ timer ./poisson -da_refine 6
on 513 x 513 grid:  error |u-uexact|_inf = 1.23567e-06
real 17.80
$ timer ./poisson -da_refine 6 -ksp_type cg -pc_type icc
on 513 x 513 grid:  error |u-uexact|_inf = 1.96296e-07
real 2.87
```

preconditioned CG: actually measure what works!

```
$ for PC in none jacobi icc cholesky; do \  
    timer ./poisson -da_refine 6 -ksp_type cg -pc_type $PC \  
    -ksp_converged_reason -ksp_compute_singularvalues; done  
Linear solve converged due to CONVERGED_RTOL iterations 1227  
Iteratively computed extreme singular values:  
    max 7.99992 min 7.52989e-05 max/min 106242.  
on 513 x 513 grid: error |u-uexact|_inf = 1.94253e-07  
real 4.91  
Linear solve converged due to CONVERGED_RTOL iterations 1227  
Iteratively computed extreme singular values:  
    max 1.99998 min 1.88247e-05 max/min 106242.  
on 513 x 513 grid: error |u-uexact|_inf = 1.94253e-07  
real 5.20  
Linear solve converged due to CONVERGED_RTOL iterations 357  
Iteratively computed extreme singular values:  
    max 1.20706 min 0.000128532 max/min 9391.18  
on 513 x 513 grid: error |u-uexact|_inf = 1.96296e-07  
real 2.83  
Linear solve converged due to CONVERGED_RTOL iterations 1  
Iteratively computed extreme singular values:  
    max 1. min 1. max/min 1.  
on 513 x 513 grid: error |u-uexact|_inf = 1.92073e-07  
real 112.32
```


preconditioned CG: the hope

- ▶ preconditioned CG was the great hope of about 1970
- ▶ *good*: preconditioning by `icc` *does* reduce iteration count by reducing condition number, relative to no preconditioning or `jacobi`
- ▶ *bad*: preconditioning by `icc` *does not* slow down the growth condition number $\kappa_2(A)$ as the grid is refined

preconditioned CG: the scaling flaw

- ▶ do with PC=none, icc:

```
$ for K in 0 1 2 3 4 5 6; do \  
  ./poisson -da_refine $K -ksp_converged_reason \  
  -ksp_type cg -pc_type PC; done
```



- ▶ we need genuinely-better preconditioning!