

Nonlinear PDEs and Jacobian options

1D example solved by 21st century means

Ed Bueler

Dept. of Mathematics and Statistics, UAF

29 March 2016

outline for today

Chapter 4 of book:

- ▶ recall 1D diffusion-reaction example `c/ch4/reaction.c`:

$$-u'' - R(u) = f(x)$$

- ▶ show:
 - evidence for convergence
 - finite-difference Jacobian
 - ▶ optionally by graph coloring
 - matrix-free Newton-Krylov
 - ▶ optionally with a preconditioning matrix

recall: problem and solution method

- ▶ balance of diffusion/reaction/source processes:

$$-u'' - R(u) = f(x), \quad u(0) = \alpha, \quad u(1) = \beta$$

- an ODE 2-point BVP, but acts like elliptic PDE (e.g. Poisson)
 - $R(u)$ is any nonlinear function
 - exact solution known when $R(u) = -\rho\sqrt{u}$ and $f(x) = 0$
- ▶ discretize on N -point structured grid by finite differences:

$$-\frac{u_{i+1} - 2u_i + u_{i-1}}{h^2} - R(u_i) = f(x_i)$$

- ▶ becomes “residual=0” equations in \mathbb{R}^N : $\mathbf{F}(\mathbf{x}) = 0$
- ▶ Newton’s method: choose \mathbf{x}_0 and iterate

$$\mathbf{J}_{\mathbf{F}}(\mathbf{x}_k) \mathbf{s} \stackrel{*}{=} -\mathbf{F}(\mathbf{x}_k), \quad \mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{s}$$

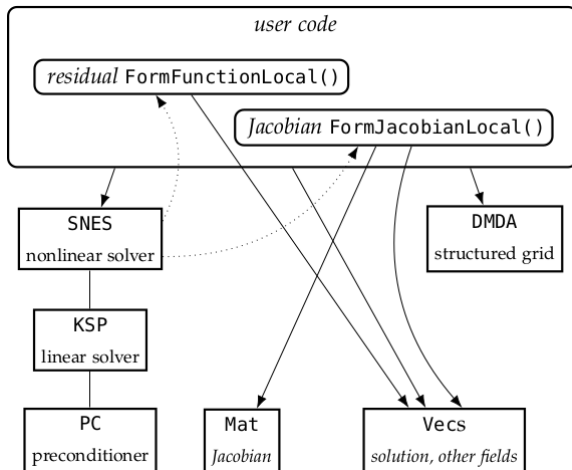
- solve each linear system $*$ by preconditioned Krylov: “Newton-Krylov”

recall: PETSc implementation

- ▶ `c/ch4/reaction.c`
- ▶ DMDA manages grid (in parallel)
- ▶ SNES manages Newton's method (in parallel)
 - has Krylov solver `KSP` and preconditioner `PC` inside
- ▶ we write these parts of code:
 - initial iterate function
 - residual function $\mathbf{F}(\mathbf{x})$, a SNES call-back
 - Jacobian $\mathbf{J}_{\mathbf{F}}(\mathbf{x})$, a SNES call-back
 - `main()`:
 - ▶ call `Create/SetFromOptions` on objects
 - ▶ call initial iterate function
 - ▶ call `SNESolve()`
 - ▶ measure error $\|\mathbf{x}_k - \mathbf{x}\|$ (relative to exact solution \mathbf{x})
 - ▶ call `Destroy` on objects



recall: structure of `c/ch4/reaction.c`

- ▶ solid arrows mean “user code acts directly on”
- ▶ dotted arrows are call-backs



basic runs


```
$ ./reaction -da_refine 4 # refine by 2^4
$ ./reaction -snes_monitor -ksp_converged_reason # info on run
$ ./reaction -snes_view # expose solvers
$ ./reaction -mat_view draw -draw_pause 1 # show matrix (L)
$ ./reaction -da_refine 4 -snes_monitor \ # show iterates (R)
    -snes_monitor_solution draw -draw_pause -1
```



```
$ mpiexec -n 4 ./reaction -da_refine 12 \ # parallel
    -snes_monitor -ksp_converged_reason
```



convergence

```
$ for N in 0 2 4 6 8 10 12 14 16; do  
> ./reaction -da_refine $N -snes_rtol 1.0e-10; done
```



distinguish meanings of “convergence”!

- ▶ previous: **convergence of numerical soln to exact soln**
 - usually the main goal
- ▶ left: **convergence of residual in the nonlinear equations**
 - Newton “outer” iteration (SNES residual)
- ▶ right: **convergence of residual in linear eqns per Newton step**
 - Krylov “inner” iteration (KSP residual)



```
$ ./reaction -snes_monitor -ksp_type preonly -pc_type lu  
$ ... -ksp_monitor -ksp_type gmres -pc_type jacobi
```


approximate Jacobians

- ▶ in our code we do have `FormJacobianLocal()`
 - ...but suppose we did not ...
implementing Jacobians can be an error-prone pain
 - of course, we *must* have a `FormFunctionLocal()` which computes $\mathbf{F}(\mathbf{x})$
- ▶ one idea for approximate Jacobian is `-snes_fd`:

$$J_{ij} = \frac{\partial F_i}{\partial x_j} \approx \frac{F_i(\mathbf{x} + \delta \mathbf{e}_j) - F_i(\mathbf{x})}{\delta}$$

- $\delta = \sqrt{\epsilon}$ where $\epsilon = 2.2 \times 10^{-16}$ is machine precision
- ▶ different idea, suitable for Krylov iteration, is `-snes_mf`:

$$\mathbf{J}\mathbf{v} \approx \frac{\mathbf{F}(\mathbf{x} + \delta \mathbf{v}) - \mathbf{F}(\mathbf{x})}{\delta}$$

approximate Jacobians: the caveat

- ▶ **neither `-snes_fd` nor `-snes_mf` scale well!**
- ▶ for different reasons:
 - `-snes_fd` does not scale well because if grid has N points then $N + 1$ residual evaluations are needed per Newton step
 - need to avoid redundant residual evaluations
 - `-snes_mf` does not scale well because only $J^k \mathbf{v}$ can be calculated, so Krylov method will not converge fast if spectrum of J is spread out
 - need to add a preconditioner
- ▶ fortunately, modified versions `-snes_fd_color` and `-snes_mf_operator` exist

finite-difference Jacobians

- ▶ options `-snes_fd` and `-snes_fd_color` only need residual implementation (`FormFunctionLocal()`)
- ▶ recall we can see how many times the residual function `FormFunctionLocal()` is evaluated:

```
$ ... -log_view |grep SNESFunctionEval
```

- ▶ for example:


```
$ ./reaction -da_refine 10
```

```
$ ./reaction -da_refine 10 -snes_fd
```

```
$ ./reaction -da_refine 10 -snes_fd_color
```

gives 4, 16390, 13 evaluations, respectively

how `-snes_fd_color` works



- ▶ graph has edge if two nodes are connected in stencil
- ▶ Jacobian J built by evaluating using columns of one color
 - graph needs $c = 3$ colors in this case
- ▶ number of `FormFunctionLocal()` evals:
 - $= q(c + 1) + 1$ if q Newton steps
 - independent of refinement level

matrix-free Newton-Krylov (“JFNK”)

- ▶ option `-snes_mf` only needs residual implementation (`FormFunctionLocal()`)
- ▶ option `-snes_mf_operator` also needs Jacobian implementation (`FormJaobianLocal()`), but it can be—and normally is—only an approximation of the Jacobian
 - $P \approx J$ will be used for preconditioning

- ▶ to illustrate the problem:

```
$ ./reaction -da_refine 10 -snes_mf
```

does not converge, while

```
$ ./reaction -da_refine 4
```

```
$ ./reaction -da_refine 4 -snes_mf
```

```
$ ./reaction -da_refine 4 -snes_fd
```

```
$ ./reaction -da_refine 4 -snes_fd_color
```

gives 4, 3225, 394, 13 evaluations, respectively

Jacobian options comparisons

	<u>no option</u>	<u>-snes_fd -snes_fd_color</u>	<u>-snes_mf</u>	<u>-snes_mf_operator</u>
only F	error	✓	✓	error
F and P	✓	✓	✓	✓
F and J	✓	✓	✓	✓

Table 4.2: Jacobian options when using SNES. Rows give which mathematical objects have been provided through user code. Symbol " P " denotes an approximate Jacobian while " J " denotes the actual Jacobian. A large check mark shows recommended usage.

	<u>no option</u>	<u>-snes_fd -snes_fd_color</u>	<u>-snes_mf -snes_mf_operator</u>
residual F	$q + 1$	$q(N + 1) + 1$ $q(c + 1) + 1$	qm_1 qm_2
Jacobian J or P	q	0 0	0 q

Table 4.3: Jacobian options compared by number of residual and Jacobian evaluations. Here q is the number of Newton iterations, N is the dimension of the problem, c is the number of colors on $G(J)$, m_1 is the dimension of the Krylov space for J , and m_2 is the Krylov space dimension for the preconditioned operator $M^{-1}J$.

► main idea:

*you have options, after implementing the residual **F(x)**, for solving without an analytical Jacobian, or with an implemented approximate Jacobian*

trying out JFNK

- ▶ for example, code `c/ch4/reaction.c` has option `-rct_noRinJ` which removes “ $R(u)$ ” part of (correct) Jacobian

- ▶ with analytical Jacobian

```
$ ./reaction -snes_monitor -da_refine 10
```

- ▶ with inexact Jacobian

```
$ ./reaction -snes_monitor -da_refine 10 \  
-rct_noRinJ
```

- ▶ with inexact Jacobian *used as preconditioner*

```
$ ./reaction -snes_monitor -da_refine 10 \  
-rct_noRinJ -snes_mf_operator
```